

This is a Python notebook with the purpose of running classification algorithms on a day trading stock data set to compare it's performance with R classification.

## Initial Data Exploration and Cleaning

Check out the Kaggle link for more information on the data set and its variables:

<https://www.kaggle.com/dawerty/cleaned-daytrading-training-data>

```
In [1]: import pandas as pd
        #load data
        df = pd.read_csv('stock.csv', header=0)
```

```
In [2]: #exploration function 1
        df.info()
```

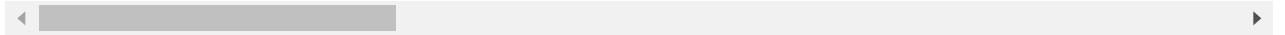
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 258729 entries, 0 to 258728
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   is_profit                             258729 non-null bool
1   sym                                   258729 non-null object
2   datetime                             258729 non-null object
3   rsi14                                258729 non-null float64
4   sma9_var                             258729 non-null float64
5   sma180_var                           258729 non-null float64
6   vwap_var                             258729 non-null float64
7   spread14_e                           258729 non-null float64
8   volume14_34_var                      258729 non-null float64
9   prev_close_var                       258729 non-null float64
10  prev_floor_var                       258729 non-null float64
11  prev_ceil_var                        258729 non-null float64
12  prev1_candle_score                   258729 non-null float64
13  prev2_candle_score                   258729 non-null float64
14  prev3_candle_score                   258729 non-null float64
15  mins_from_start                     258729 non-null float64
16  valley_interval_mins                 258729 non-null float64
17  valley_close_score                   258729 non-null float64
18  valley_rsi_score                     258729 non-null float64
19  day_open_var                        258729 non-null float64
20  open_from_prev_close_var             258729 non-null float64
21  ceil_var                            258729 non-null float64
22  floor_var                            258729 non-null float64
dtypes: bool(1), float64(20), object(2)
memory usage: 43.7+ MB
```

```
In [3]: #exploration function 2
        df.describe()
```

```
Out[3]:
```

	rsi14	sma9_var	sma180_var	vwap_var	spread14_e	volume14_34_var	p
<b>count</b>	258729.000000	258729.000000	258729.000000	258729.000000	258729.000000	258729.000000	2
<b>mean</b>	34.566266	-0.002659	-0.011785	-0.009576	0.000885	-0.043627	
<b>std</b>	5.463582	0.002344	0.014630	0.010056	0.000765	0.310340	

	rsi14	sma9_var	sma180_var	vwap_var	spread14_e	volume14_34_var	p
min	6.140843	-0.019984	-0.199053	-0.126805	0.000002	-1.000000	
25%	31.112562	-0.003322	-0.016985	-0.013535	0.000406	-0.230678	
50%	34.838873	-0.001950	-0.008588	-0.007803	0.000653	-0.047021	
75%	38.210769	-0.001157	-0.003302	-0.003722	0.001089	0.134563	
max	71.815499	0.002438	0.183595	0.071907	0.012104	1.428571	



```
In [4]: #cleaning: changing is_profit data types
df.is_profit = df.is_profit.astype('category').cat.codes
del df["datetime"] #drop datetime column
del df["sym"] #drop sym column
```

```
In [5]: #exploration function 3
df.head
```

```
Out[5]: <bound method NDFrame.head of
ar spread14_e \
0      1  30.509761 -0.006223 -0.022679 -0.017526  0.000620
1      1  46.452741 -0.001062 -0.004721 -0.007713  0.000695
2      1  34.336224 -0.004443 -0.016648 -0.016589  0.000518
3      0  36.584676 -0.001006  0.005697 -0.004279  0.000327
4      1  29.113480 -0.000950  0.002626 -0.001767  0.000286
...
258724 0  32.602899 -0.002293 -0.033872 -0.011582  0.001114
258725 0  37.355860 -0.002370 -0.021928 -0.015194  0.000470
258726 1  41.550637 -0.001991 -0.014774 -0.003803  0.000993
258727 0  35.433061 -0.005427 -0.012600 -0.015767  0.001918
258728 1  37.648564 -0.004007 -0.016134 -0.016528  0.002374

      volume14_34_var prev_close_var prev_floor_var prev_ceil_var ... \
0      -0.006472      -0.037037      -0.012658      -0.047328 ...
1      0.280249      -0.031893      -0.007384      -0.042239 ...
2      0.284800      0.011396      0.023360      -0.014706 ...
3     -0.514448      0.017371      0.023517      0.001266 ...
4     -0.033291      0.019482      0.035230      -0.001628 ...
...
258724  0.462271      -0.029817      -0.028145      -0.060000 ...
258725 -0.238875      -0.038417      -0.036760      -0.068333 ...
258726 -0.062510      0.057895      0.105147      0.050330 ...
258727  0.046143      0.004747      0.015770      -0.095270 ...
258728 -0.000039      0.002035      0.013027      -0.097713 ...

      prev2_candle_score prev3_candle_score mins_from_start \
0      0.000000      0.000000      103.0
1      0.001062      0.000504      265.0
2     -0.001020      0.000000      278.0
3     -0.000210      0.000000      110.0
4     -0.000012      0.000000      229.0
...
258724  0.000000      -0.000590      85.0
258725  0.001488      0.000209      247.0
258726  0.000000      0.003775      285.0
258727 -0.001341      0.002687      146.0
258728 -0.001509      -0.000506      216.0

      valley_interval_mins valley_close_score valley_rsi_score \
0      50.0      0.425532      0.758046
```

1	67.0	0.633584	10.958588
2	13.0	0.306356	2.964667
3	8.0	0.042142	2.599359
4	29.0	0.224383	0.091923
...	...	...	...
258724	19.0	0.118066	7.540690
258725	46.0	0.563293	10.917935
258726	20.0	0.077993	1.535276
258727	12.0	0.235696	4.331218
258728	9.0	0.002033	10.013797

	day_open_var	open_from_prev_close_var	ceil_var	floor_var
0	-0.032058	-0.005144	-0.034554	0.000802
1	-0.026887	-0.005144	-0.029397	0.006148
2	-0.003935	0.015391	-0.030638	0.000000
3	0.003596	0.013725	-0.014746	0.003596
4	0.008738	0.010651	-0.016754	0.011470
...	...	...	...	...
258724	-0.021400	-0.008601	-0.030373	0.000000
258725	-0.030075	-0.008601	-0.038968	0.000000
258726	0.070815	-0.012065	-0.037562	0.070815
258727	0.003386	0.001356	-0.033909	0.003386
258728	0.000677	0.001356	-0.036518	0.000677

[258729 rows x 21 columns]>

In [6]: `#exploration function 3 (cont.)`  
`df.tail`

Out[6]: <bound method NDFrame.tail of

ar	spread14_e	\	is_profit	rsi14	sma9_var	sma180_var	vwap_v
0	1	30.509761	-0.006223	-0.022679	-0.017526	0.000620	
1	1	46.452741	-0.001062	-0.004721	-0.007713	0.000695	
2	1	34.336224	-0.004443	-0.016648	-0.016589	0.000518	
3	0	36.584676	-0.001006	0.005697	-0.004279	0.000327	
4	1	29.113480	-0.000950	0.002626	-0.001767	0.000286	
...	...	...	...	...	...	...	
258724	0	32.602899	-0.002293	-0.033872	-0.011582	0.001114	
258725	0	37.355860	-0.002370	-0.021928	-0.015194	0.000470	
258726	1	41.550637	-0.001991	-0.014774	-0.003803	0.000993	
258727	0	35.433061	-0.005427	-0.012600	-0.015767	0.001918	
258728	1	37.648564	-0.004007	-0.016134	-0.016528	0.002374	

	volume14_34_var	prev_close_var	prev_floor_var	prev_ceil_var	...	\
0	-0.006472	-0.037037	-0.012658	-0.047328	...	
1	0.280249	-0.031893	-0.007384	-0.042239	...	
2	0.284800	0.011396	0.023360	-0.014706	...	
3	-0.514448	0.017371	0.023517	0.001266	...	
4	-0.033291	0.019482	0.035230	-0.001628	...	
...	...	...	...	...	...	
258724	0.462271	-0.029817	-0.028145	-0.060000	...	
258725	-0.238875	-0.038417	-0.036760	-0.068333	...	
258726	-0.062510	0.057895	0.105147	0.050330	...	
258727	0.046143	0.004747	0.015770	-0.095270	...	
258728	-0.000039	0.002035	0.013027	-0.097713	...	

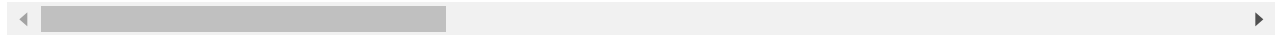
  

	prev2_candle_score	prev3_candle_score	mins_from_start	\
0	0.000000	0.000000	103.0	
1	0.001062	0.000504	265.0	
2	-0.001020	0.000000	278.0	
3	-0.000210	0.000000	110.0	
4	-0.000012	0.000000	229.0	
...	...	...	...	
258724	0.000000	-0.000590	85.0	
258725	0.001488	0.000209	247.0	



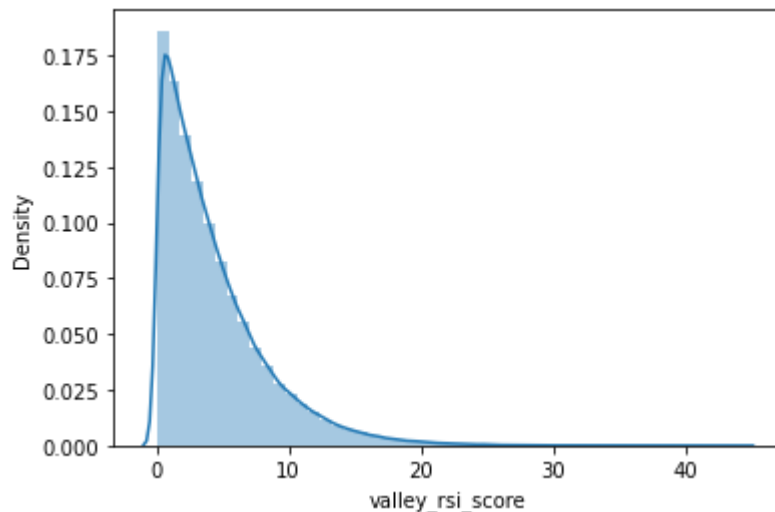
	is_profit	rsi14	sma9_var	sma180_var	vwap_var	spread14_e	volume14_34_var	prev_close_var
...	...	...	...	...	...	...	...	...
258724	False	False	False	False	False	False	False	False
258725	False	False	False	False	False	False	False	False
258726	False	False	False	False	False	False	False	False
258727	False	False	False	False	False	False	False	False
258728	False	False	False	False	False	False	False	False

258729 rows × 21 columns



```
In [9]: #exploration graph 1
from plotnine import *
import seaborn as sns
sns.distplot(df['valley_rsi_score']);
```

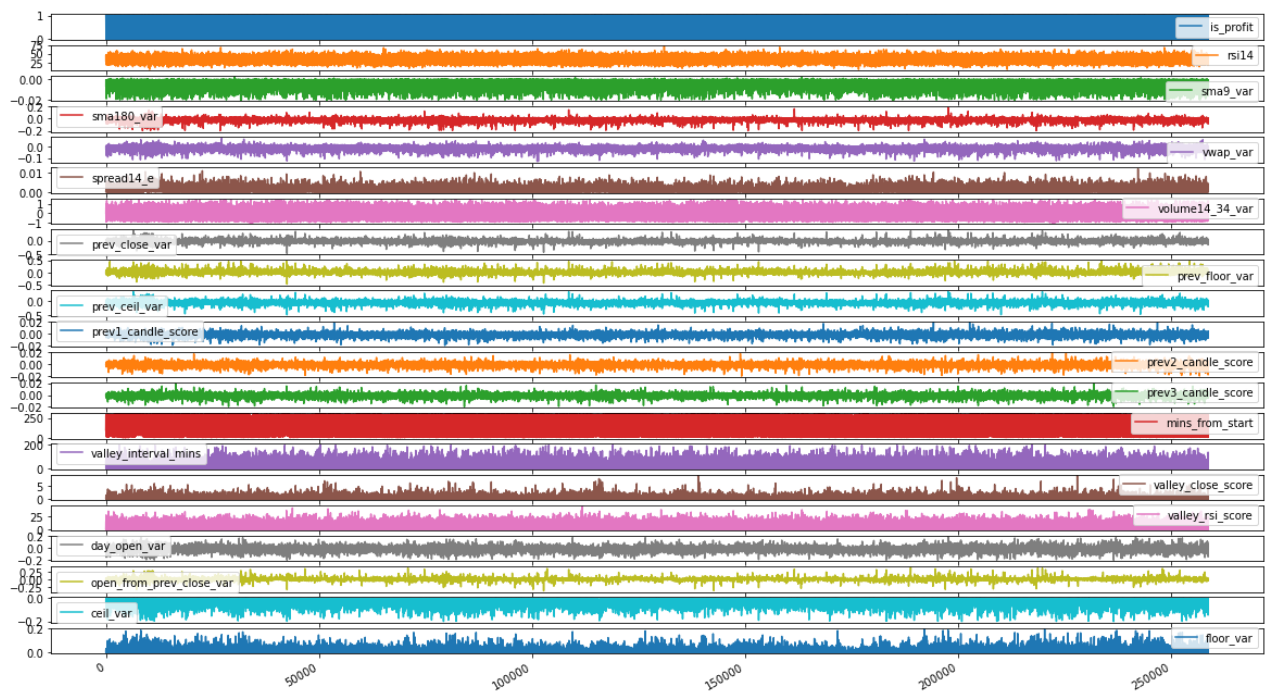
C:\Users\RaxyR\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [37]: #exploration graph 2
import matplotlib.pyplot as plt
plt.figure(figsize=(16,10), dpi= 80)
sns.kdeplot(df.loc[df['is_profit'] == 0, "volume14_34_var"], shade=True, color="deeppin
sns.kdeplot(df.loc[df['is_profit'] == 1, "volume14_34_var"], shade=True, color="dodgerb
```

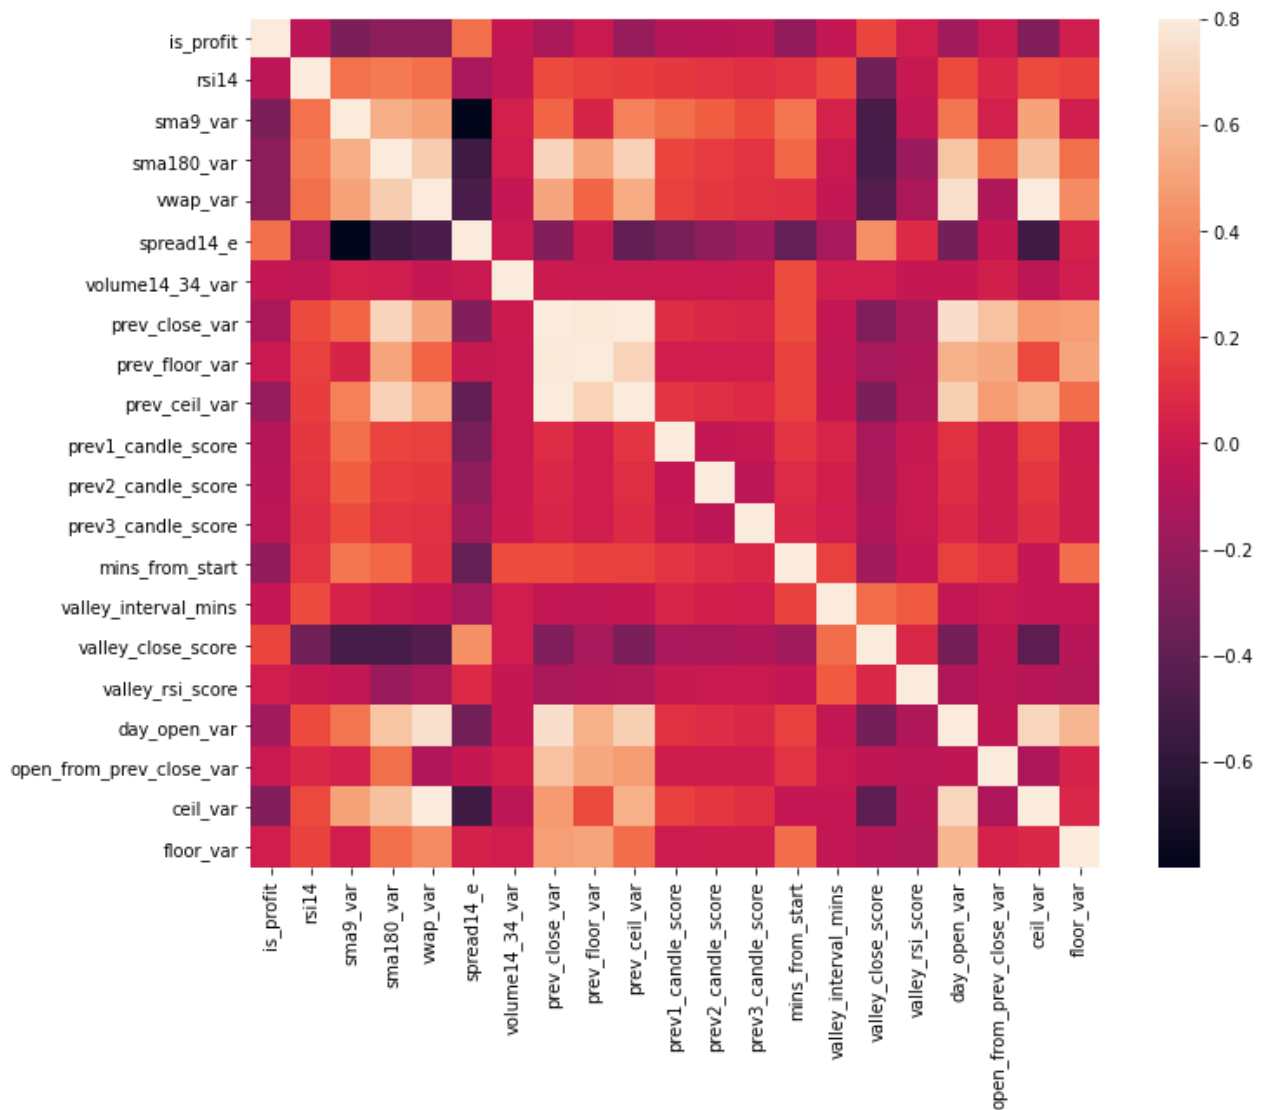
Out[37]: <AxesSubplot:xlabel='volume14\_34\_var', ylabel='Density'>





```
In [26]: #exploration graph 5
corrmat = df.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, vmax=.8, square=True)
```

```
Out[26]: <AxesSubplot:>
```



## Modeling

```
In [11]: # train test split
from sklearn.model_selection import train_test_split
X = df.iloc[:, df.columns != 'is_profit']
y = df.iloc[:, df.columns == 'is_profit']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state=1234)
print('train size:', X_train.shape)
print('test size:', X_test.shape)
```

train size: (194046, 20)  
test size: (64683, 20)

```
In [13]: #logistic regression model
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(solver='lbfgs', max_iter=10000)
clf.fit(X_train, y_train)
clf.score(X_train, y_train)

# make predictions
pred1 = clf.predict(X_test)
from sklearn.metrics import classification_report
print(classification_report(y_test, pred1))
```



```
probas_pred = clf.predict_proba(X_test)[:,-1]
```

C:\Users\RaxyR\anaconda3\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning:

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

	precision	recall	f1-score	support
0	0.70	0.88	0.78	40206
1	0.66	0.38	0.48	24477
accuracy			0.69	64683
macro avg	0.68	0.63	0.63	64683
weighted avg	0.68	0.69	0.67	64683

```
In [32]: #naive bayes model
from sklearn.naive_bayes import GaussianNB #note: using GaussianNB due to negative numb
clf = GaussianNB()
clf.fit(X_train, y_train)
clf.score(X_train, y_train)

# make predictions
pred2 = clf.predict(X_test)
print(classification_report(y_test, pred2))
```

	precision	recall	f1-score	support
0	0.68	0.88	0.77	40206
1	0.62	0.34	0.44	24477
accuracy			0.67	64683
macro avg	0.65	0.61	0.60	64683
weighted avg	0.66	0.67	0.64	64683

C:\Users\RaxyR\anaconda3\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning:

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
In [33]: #decision tree model
from sklearn.tree import DecisionTreeClassifier
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
clf.score(X_train, y_train)

# make predictions
pred3 = clf.predict(X_test)
print(classification_report(y_test, pred3))
```

	precision	recall	f1-score	support
0	0.69	0.68	0.69	40206
1	0.49	0.50	0.49	24477
accuracy			0.61	64683
macro avg	0.59	0.59	0.59	64683

weighted avg	0.61	0.61	0.61	64683
--------------	------	------	------	-------

Logistic regression actually performed just slightly worse with Python than with R (accuracy .69 < .70). Naive bayes also performed slightly worse with Python (.67 < .68). Finally, the decision tree performed slightly better with Python (.69 > .68). Ranking these three algorithms in Python, we have decision tree and logistic regression tied and then naive bayes. This is very similar with the R project with the contention between the decision tree and logistic regression being with the margins of 1-2% accuracy.

Some of the predictors may not have been independent so the naive assumption that they are may have limited the performance of the naive bayes model. This is most likely the reason it was outperformed by logistic regression and the decision tree. Logistic regression searches for a single linear decision boundary whereas the decision tree partitions the feature space into half spaces for a boundary but in this case the effect was more or less the same. However, because decision trees are so flexible, the model may have been prone to overfitting and logistic regression was less susceptible here. Maybe if any pruning was done, the accuracy could have increased. All in all, this was a battle of bias-variance tradeoff and logistic won, very slightly, and naive bayes struggled against the size of the data set.

Interestingly, naive bayes ran the fastest in Python and logistic regression the slowest, essentially the opposite of what occurred in R.

Personally, I lean towards machine learning in R rather than Python perhaps because I have had a longer history of experience with it but also because of ease of explanatory data analysis. There are a wide array of statistical functions and options with plotting and even though I was able to replicate ggplot2 in Python with the plotnine library, there are less nuances with R. I do enjoy the power of Python in that it can be used for almost anything being very flexible. R being very data analysis oriented makes it very useful for just focusing on that. So, to summarize, I believe data analysis and visualization is better done in R but more complicated modeling and deep learning should be done with Python. I think in the long run, I would use Python as it is more powerful and versatile but R will be a personal favorite with it's ease of use.